

YIXIAO(CLAIRE) LING

(+001) 267 206 2692

liing@seas.upenn.edu; <https://github.com/yixiao-ling?tab=repositories>; <https://www.linkedin.com/in/yixiao-ling/>

EDUCATION

University of Pennsylvania | Penn Engineering

Aug. 2023 – May. 2025

M.S.E. in Data Science

Relevant Courses: Data Structure and Algorithms, Machine Learning, Distributed Systems, Deep Learning, Large Language Models, Computer Vision

GPA: 3.96/4.00

University of Nottingham | School of Mathematical Science

Sep. 2019-Jul. 2023

B.S. in Mathematics and Applied Mathematics

Relevant Courses: Statistical Models and Methods, Probability Models and Methods, Econometrics

GPA: 3.92/4.00 (Top 5%)

Honors: Dean's scholarship

The London School of Economics and Political Science

Jun. 2021 – Jul. 2021

Summer Course Program in Computational Methods for Financial Mathematics

Grade: A+

TECHNICAL SKILLS

ML & AI Engineering Skills: Experience in end-to-end ML and implementing CI/CD pipelines. Experience in deploying models to cloud services (AWS, Azure). Experienced in working with ML platforms (Dataiku, SageMaker). Expertise in LLM technologies and AI infra. Strong experience with Generative AI models.

Data Science Skills: Proficiency in large-scale systems data analysis and statistical modelling.

Programming Languages: C++, Python, SQL, Java, Spark, MATLAB

Platforms & Tools: Google Cloud, Amazon Web Services, RDMS, MYSQL, R Studio, Tableau, Power BI

WORK EXPERIENCE

AI Engineer, SAP

Jun.2025 – Now

- Developed **full-stack distributed multi-modal AI system** (chatbot and **slides auto-generation**) for sales professionals, enabling rapid customer analysis with high precision and actionable insights.
- Engineered real-time data ingestion from multiple sources, including document processing and web scraping, incorporating **automated quality validation**, **S3 cloud storage**, and intelligent **metadata management** for seamless content extraction.
- Integrate multiple LLM providers (**OpenAI, Gemini, Claude**), **optimized algorithms** and built **caching** layer with **Redis** for session management, file uploads and API response caching.
- Developed **concurrency management** using **ThreadPool**, **asyncio**, **rate limiting**, Redis streams and **process tools** for CPU-intensive tasks, scaling system to handle 10,000+ concurrent requests with **queue management**.
- Established **monitoring and alerting** with custom metrics for API response times, error rates and resource utilization.

Built **Analytics dashboard** with real-time metrics, usage tracking, and performance monitoring for business teams.

AI Engineer Intern, Alibaba Cloud

Feb. 2024 – Sep.2024

- Designed and built a **scalable RAG-based chatbot system** on distributed architecture using **cloud service**, deployed with **FastAPI** and maintained system scalability & performance.
- Developed a **multimodal module**, optimized **LLM reasoning** by **fine-tuning**, **prompting** techniques and developed agents to improve conversation quality.
- Developed ElasticSearch (ES) **vector database**, optimized database loader, semantic **search** and **ranking algorithms**, increasing the recall accuracy of 4.5% and recall speed of 15%.
- Used the **CI/CD pipeline** to automate deploying the optimized LLM model to production and created a diagnostics platform for real-time monitoring & analysis of LLM **distributed training**, ensuring effective **performance tracking**.

Generative AI Intern, Wharton Analytics

Sep. 2023 – Feb.2024

- Developed a **Machine Learning** solution for Hearst Corporation to automatically align magazine content with marketplace taxonomy.
- Employed and optimized **PySpark** for **distributed data processing** on large volumes of HTML data. Leveraged Delta Lake to manage data versioning and maintain a reliable, auditable data pipeline.
- Used **AWS** for building, training, and deploying the NLP models. **Fine-tuned BERT** for NLP semantic analysis.
- Achieved a 78% tagging accuracy rate, leading to a 15% increase in user engagement and a 10% boost in sales.

Machine Learning Engineer Intern, Experian

Jan. 2023-Sep. 2023

- Developed data preprocessing pipeline on 13GB user interaction data to prepare for building risk control model.
- Conducted **feature engineering**, selecting 35 key features from 5000+ features by feature importance and variable correlation. Derived 89 new variables based on real business.
- Deployed **risk control model** to determine whether to give credits to clients by using machine learning models, such as **logistic regression**, **XGBoost** and **LightGBM**, achieving accuracy of 0.85 classifying good & bad clients.